

HW2 - Tardella - part I: exercises 1, 2 and 3

Your Last+First Name: Rossini Valerio, Your Matricola: 1613638

Part 1

Dugongs - part I

1a answer)

We are doing a nonconjugate Bayesian analysis of a data set that contains information about dugongs. The data are length (the dependent variable) and age (the independent variable) measurements for 27 (i.e. we have $n = 27$) captured dugongs, and they are modelled with the use of a nonlinear growth curve with no inflection point and an asymptote as X_i tends to infinity:

$$Y_i \sim \text{Normal}(\mu_i, \tau) \quad \text{with } i = 1, \dots, 27$$

where $\mu_i = \mathbb{E}[Y_i | X_i = x_i] = f(x_i) = \alpha - \beta \gamma^{x_i}$

The model parameters are $\alpha \in (1, \infty)$, $\beta \in (1, \infty)$, $\gamma \in (0, 1)$ and $\tau^2 \in (0, \infty)$ and we consider the following prior distribution:


$$\alpha \sim N(0, \sigma_\alpha^2)$$

$$\beta \sim N(0, \sigma_\beta^2)$$

$$\gamma \sim \text{Unif}(0, 1)$$

$$\tau^2 \sim \text{IG}(a, b)$$

The graph is given below:



After that we load the data set, we can see graphically a simple relation between the two variables

```
df=read.csv("dugongs.txt",sep="\t")

# let's see the variables into the dataframe
names(df)
```


```
## [1] "Age"      "Length"
# number of dugongs present in the dataframe
n = dim(df)[1]
n

## [1] 27
# the dependent variable represents the length of the dugongs, while the dependent
# variable represents the age of the dugongs
x = df$Age
y = df$Length

## x = age
# x = c( 1.0,  1.5,  1.5,  1.5, 2.5,  4.0,  5.0,  5.0,  7.0,
#       8.0,  8.5,  9.0,  9.5, 9.5, 10.0, 12.0, 12.0, 13.0,
#       13.0, 14.5, 15.5, 15.5, 16.5, 17.0, 22.5, 29.0, 31.5)

## y = lenght
# y = c(1.80, 1.85, 1.87, 1.77, 2.02, 2.27, 2.15, 2.26, 2.47,
#       2.19, 2.26, 2.40, 2.39, 2.41, 2.50, 2.32, 2.32, 2.43,
#       2.47, 2.56, 2.65, 2.47, 2.64, 2.56, 2.70, 2.72, 2.57)

plot(x,y, xlab = "Age", ylab = "Length", main = "Non linear regression")
```



1b answer)

The likelihood function can be derived in this way:

$$\begin{aligned}
L(\alpha, \beta, \gamma, \tau | y, x) &= \prod_{i=1}^n \frac{1}{(2\pi\tau^2)^{\frac{1}{2}}} \cdot \exp \left\{ -\frac{(y_i - \mu_i)^2}{2\tau^2} \right\} \mathbb{I}_{(1,\infty)}(\alpha) \mathbb{I}_{(1,\infty)}(\beta) \mathbb{I}_{(0,1)}(\gamma) \mathbb{I}_{(0,\infty)}(\tau^2) \\
L(\alpha, \beta, \gamma, \tau | y, x) &= \frac{1}{(2\pi\tau^2)^{\frac{n}{2}}} \cdot \exp \left\{ -\frac{1}{2\tau^2} \sum_{i=1}^n (y_i - \mu_i)^2 \right\} \mathbb{I}_{(1,\infty)}(\alpha) \mathbb{I}_{(1,\infty)}(\beta) \mathbb{I}_{(0,1)}(\gamma) \mathbb{I}_{(0,\infty)}(\tau^2) \\
&= \frac{1}{(2\pi\tau^2)^{\frac{n}{2}}} \cdot \exp \left\{ -\frac{1}{2\tau^2} \sum_{i=1}^n (y_i - \alpha + \beta\gamma^{x_i})^2 \right\} \mathbb{I}_{(1,\infty)}(\alpha) \mathbb{I}_{(1,\infty)}(\beta) \mathbb{I}_{(0,1)}(\gamma) \mathbb{I}_{(0,\infty)}(\tau^2)
\end{aligned}$$

1c answer)

Before writing the expression of the joint prior distribution of the parameters, we need to compute separately the prior of each parameter:

$$\begin{aligned}
\pi(\alpha) &= \frac{1}{\sqrt{2\pi\sigma_\alpha^2}} \exp \left\{ -\frac{\alpha^2}{2\sigma_\alpha^2} \right\} \mathbb{I}_{(1,\infty)}(\alpha) \propto \exp \left\{ -\frac{\alpha^2}{2\sigma_\alpha^2} \right\} \mathbb{I}_{(1,\infty)}(\alpha) \\
\pi(\beta) &= \frac{1}{\sqrt{2\pi\sigma_\beta^2}} \exp \left\{ -\frac{\beta^2}{2\sigma_\beta^2} \right\} \mathbb{I}_{(1,\infty)}(\beta) \propto \exp \left\{ -\frac{\beta^2}{2\sigma_\beta^2} \right\} \mathbb{I}_{(1,\infty)}(\beta) \\
\pi(\gamma) &= \mathbb{I}_{(0,1)}(\gamma) \\
\pi(\tau^2) &= \frac{b^a}{\Gamma(a)} \tau^{2(-a-1)} \exp \left\{ -\frac{b}{\tau^2} \right\} \mathbb{I}_{(0,\infty)}(\tau^2) \propto \tau^{2(-a-1)} \exp \left\{ -\frac{b}{\tau^2} \right\} \mathbb{I}_{(0,\infty)}(\tau^2)
\end{aligned}$$

The joint prior can be obtained as:

$$\begin{aligned}
\pi(\theta) = \pi(\alpha)\pi(\beta)\pi(\gamma)\pi(\tau^2) &= \frac{1}{\sqrt{2\pi\sigma_\alpha^2}} \exp \left\{ -\frac{\alpha^2}{2\sigma_\alpha^2} \right\} \cdot \frac{1}{\sqrt{2\pi\sigma_\beta^2}} \exp \left\{ -\frac{\beta^2}{2\sigma_\beta^2} \right\} \cdot \mathbb{I}_{(0,1)}(\gamma) \cdot \frac{b^a}{\Gamma(a)} \tau^{2(-a-1)} \exp \left\{ -\frac{b}{\tau^2} \right\} \\
&\propto \exp \left\{ -\frac{\alpha^2}{2\sigma_\alpha^2} \right\} \cdot \exp \left\{ -\frac{\beta^2}{2\sigma_\beta^2} \right\} \cdot \mathbb{I}_{(0,1)}(\gamma) \cdot \tau^{2(-a-1)} \exp \left\{ -\frac{b}{\tau^2} \right\}
\end{aligned}$$

About the values of σ_α^2 , σ_β^2 , a and b , I chose: $\sigma_\alpha^2 = 10000$, $\sigma_\beta^2 = 10000$, $a = 0.001$ and $b = 0.001$

$$\begin{aligned}
\alpha &\sim N(0, \sigma_\alpha^2) \implies \alpha \sim N(0, 10000) \\
\beta &\sim N(0, \sigma_\beta^2) \implies \beta \sim N(0, 10000) \\
\gamma &\sim \text{Unif}(0, 1) \\
\tau^2 &\sim IG(a, b) \implies \tau^2 \sim IG(0.001, 0.001)
\end{aligned}$$

1d answer)

We compute numerically the maximum likelihood estimate for the vector of parameters of interest $(\alpha, \beta, \gamma, \tau)$ and compare it with the Maximum-a-Posteriori estimate in the following R code:

```

# load the packages
library(MCMCpack)

## Warning: package 'MCMCpack' was built under R version 3.3.3
## Loading required package: coda
## Warning: package 'coda' was built under R version 3.3.3
## Loading required package: MASS

## ##
## ## Markov Chain Monte Carlo Package (MCMCpack)
## ## Copyright (C) 2003-2017 Andrew D. Martin, Kevin M. Quinn, and Jong Hee Park
## ##
## ## Support provided by the U.S. National Science Foundation
## ## (Grants SES-0350646 and SES-0350613)
## ##

# initial parameters of alpha, beta, gamma and tau (let's chose them randomly)
init_parameters = c(1.5, 1.5, 0.5, 1)

# compute the log-likelihood function
log_L = function(pars) {
  alpha = pars[1]
  beta = pars[2]
  gamma = pars[3]
  tau = pars[4]
  mu = alpha - beta*(gamma^(x))
  #I = *(alpha>1)*(beta>1)*(gamma>0 & gamma<1)*(tau>0)
  out = sum(dnorm(y, mu, tau, log=T)) #*I
  return(out)
}

# find maximum likelihood estimates for the vector of parameters of interest
opt_parameters = optim(par = init_parameters, log_L, control = list(fnscale=-1))$par

## Warning in dnorm(y, mu, tau, log = T): Si è prodotto un NaN


# the maximum likelihood estimates for the vector of
# parameters of interest are respectively:
opt_parameters

## [1] 2.66662195 0.97255678 0.87349199 0.08316385

# plot the data and let's see graphically the result of the non linear regression if we
# chose the parameters alpha, beta, gamma and tau with the command "optim"
plot(x,y, xlab = "Age", ylab = "Length",
     main = "Non linear regression with MLE parameters")
lines(x, opt_parameters[1]-opt_parameters[2]*opt_parameters[3]^x, type = "l",
     col = "dark red", lwd = 2.5)

```

Non linear regression with MLE parameters



```
# compute, for every parameter, the corresponding prior
prior_alpha = function(alpha) return(dnorm(alpha, 0, 10000))
prior_beta = function(beta) return(dnorm(beta, 0, 10000))
prior_gamma = function(gamma) return(dunif(gamma))
prior_tau = function(tau,a,b) return(dinvgamma(tau, 0.001, 0.001))

# the joint prior distribution can be computed as product of the prior of alpha,
# beta, gamma and tau
joint_prior = function(pars){
  alpha = pars[1]
  beta = pars[2]
  gamma = pars[3]
  tau = pars[4]
  out = prior_alpha(alpha)*prior_beta(beta)*prior_gamma(gamma)*prior_tau(tau,a,b)
  return(out)
}

# compute the log-posterior
log_posterior = function(pars){
  alpha = pars[1]
  beta = pars[2]
  gamma = pars[3]
  tau = pars[4]
  out = log_L(pars) + log(prior_alpha(alpha)) + log(prior_beta(beta)) +
    log(prior_gamma(gamma)) + log(prior_tau(tau))
  return(out)
}
```

```

}

# find Maximum-a-Posteriori estimate for the vector of parameters of interest
opt_parameters_posterior = optim(par = init_parameters, log_posterior,
                                control = list(fnscale=-1))$par


## Warning in dnorm(y, mu, tau, log = T): Si è prodotto un NaN
## Warning in log(x): Si è prodotto un NaN
## Warning in dnorm(y, mu, tau, log = T): Si è prodotto un NaN
## Warning in log(x): Si è prodotto un NaN
## Warning in dnorm(y, mu, tau, log = T): Si è prodotto un NaN
## Warning in log(x): Si è prodotto un NaN
# the Maximum-a-Posteriori estimate for the vector of parameters
# of interest are respectively:
opt_parameters_posterior

## [1] 2.66664140 0.97250181 0.87350365 0.08168616

# plot the data and let's see graphically the result of the non linear regression if we
# chose the parameters alpha, beta, gamma and tau with the command "optim"
plot(x,y, xlab = "Age", ylab = "Length",
     main = "Non linear regression with MAP parameters")
lines(x, opt_parameters_posterior[1]-opt_parameters_posterior[2]
      *opt_parameters_posterior[3]^x, type = "l", col = "red", lwd = 2.5)


```

Non linear regression with MAP parameters




```
# let's compare the last plot with the initial plot
par(mfrow = c(1,2))
plot(x,y, xlab = "Age", ylab = "Length",
     main = "Non linear regression \nwith MLE parameters")
lines(x, opt_parameters[1]-opt_parameters[2]*opt_parameters[3]^x, type = "l",
     col = "dark red", lwd = 2.5)
plot(x,y, xlab = "Age", ylab = "Length",
     main = "Non linear regression \nwith MAP parameters")
lines(x, opt_parameters_posterior[1]-opt_parameters_posterior[2]
     *opt_parameters_posterior[3]^x, type = "l", col = "red", lwd = 2.5)
```

**Non linear regression
with MLE parameters**



**Non linear regression
with MAP parameters**



```
par(mfrow = c(1,1))

# comparison between estimated MLE parameters and estimated MAP parameters
comparison = rbind("estimated MLE parameters:" = opt_parameters,
                  "estimated MAP parameters:" = opt_parameters_posterior)
colnames(comparison) = c("alpha", "beta", "gamma", "tau")
comparison
```

```
##              alpha      beta      gamma      tau
## estimated MLE parameters: 2.666622 0.9725568 0.8734920 0.08316385
## estimated MAP parameters: 2.666641 0.9725018 0.8735037 0.08168616
```

*# we can say that the maximum likelihood estimates and the Maximum-a-Posteriori estimate
for the vector of parameters of interest are very similar*

Part 2

Acceptance/Rejection method - part I

2a answer)

If we consider the A-R algorithm in the most general form, our target is known up to a multiplicative constant

$$\tilde{f}_X(x) = c \cdot f_X(x),$$

then $f_X(x) = \frac{\tilde{f}_X(x)}{c}$.

We want a probability density q and a suitable constant k such that

$$f_X(x) \leq kq(x) \quad \forall x \in \mathcal{X}$$

About the Acceptance/Rejection method, to simulate we need:

- 1) target distribution written $f_X(x)$ from which we can compute it
- 2) a candidate distribution q that is easy to simulate, so we can simulate Y_1, \dots, Y_n from candidate distribution $Y \sim q$

The candidate must be able to dominate target distribution:

$$\exists k > 0 \quad s.t. \quad f_X(x) \leq kq(x) \quad \forall x \in \mathcal{Y}$$

For each $Y \sim f_U$ we use an auxiliary experiment

$$Y = \begin{cases} Y^A & \text{if } U \leq \frac{f_X(x)}{kq(y)} \\ Y^R & \text{otherwise} \end{cases}$$

About the acceptance probability, we have that:

- 1) $Y \sim q$ with support \mathcal{Y} (suppose $\mathcal{Y} = \mathbb{R}$)
- 2) $U \sim Unif[0, 1]$
- 3) Y, U independent
- 4) If $U \leq \frac{\tilde{f}_X(y)}{ckq_Y(y)}$, we accept.

If we consider the Acceptance-Rejection algorithm in the most general form and denote with $\theta = Y^A$ the random variable obtained with the algorithm, we can compute the unconditional acceptance probability and the conditional acceptance probability

The unconditional acceptance probability is given by:

$$Pr(Y = Y^A | Y = y) = Pr\left(U \leq \frac{\tilde{f}_\Theta(Y)}{ckq_Y(Y)} | Y = y\right) = Pr\left(U \leq \frac{\tilde{f}_\Theta(y)}{ckq_Y(y)}\right) = \frac{\tilde{f}_\Theta(y)}{ckq_Y(y)}$$

The conditional acceptance probability is given by:

$$Pr(Y = Y^A) = \int_{\mathcal{Y}} Pr(Y = Y^A | Y = y) q_Y(y) dy = \int_{\mathcal{Y}} \frac{\tilde{f}_\Theta(y)}{ckq_Y(y)} q_Y(y) dy = \int_{\mathcal{Y}} \frac{\tilde{f}_\Theta(y)}{ck} dy = \frac{1}{ck}$$

2b answer)

We can prove that θ has the desired target distribution

$$\begin{aligned} Pr(Y^A < \theta) &= Pr(Y < \theta | Y = Y^A) = \frac{\int_{-\infty}^{\theta} Pr(Y = Y^A | Y = y) q_Y(y) dy}{\int_{-\infty}^{+\infty} Pr(Y = Y^A | Y = y) q_Y(y) dy} \\ &= \frac{\int_{-\infty}^{\theta} \frac{\tilde{f}_{\Theta}(y)}{ckq_Y(y)} q_Y(y) dy}{\int_{-\infty}^{+\infty} \frac{\tilde{f}_{\Theta}(y)}{ckq_Y(y)} q_Y(y) dy} = \frac{\int_{-\infty}^{\theta} \frac{\tilde{f}_{\Theta}(y)}{ck} dy}{\int_{-\infty}^{+\infty} \frac{\tilde{f}_{\Theta}(y)}{ck} dy} = \frac{\frac{1}{ck} \int_{-\infty}^{\theta} \tilde{f}_{\Theta}(y) dy}{\frac{1}{ck} \int_{-\infty}^{+\infty} \tilde{f}_{\Theta}(y) dy} = \int_{-\infty}^{\theta} \tilde{f}_{\Theta}(y) dy = F_{\Theta}(\theta) \end{aligned}$$

2c answer)

We can show how in Bayesian inference we could use simulations from the prior (auxiliary density) to get a random draw from the posterior (target distribution) without knowing the proportionality constant

Let's consider the case of Bayesian inference where our target distribution is:

$$\tilde{f}_X(x) = \pi(\theta|x) = c \cdot \pi(\theta) L_X(\theta) \implies f_X(\theta) = \frac{\tilde{f}_X(\theta)}{c} = \frac{c \cdot \pi(\theta) L_X(\theta)}{c} = \pi(\theta) L_X(\theta)$$

and the candidate $\theta \sim q(\theta) = \pi(\theta)$

The acceptance rule is:

$$U \leq \frac{f_X(\theta)}{kq(\theta)} = \frac{\pi(\theta)L_X(\theta)}{k\pi(\theta)} = \frac{L_X(\theta)}{k}$$

The optimal choice of k is obtained by maximizing the likelihood, then values are accepted if $U \leq \frac{L(\theta)}{\max_{\theta \in \Theta} L(\theta)}$

2d answer)

It's possible to find problem with A-R method, for example, when n (number of observations) is large: the likelihood is highly peaked (and so also the posterior, "precise measurement") relative to the prior, therefore the acceptance probability could be very small. In general, when the likelihood is concentrated around $\hat{\theta}_{MLE}$, it's difficult to have a good result, also in terms of MC approximation (much variability).

```
# comparison between Beta posteriors with different number of observations

alpha = 3
beta = 3

n_sim_small=20
# generation of sample from Beta with 20 observation
sample_small = rbeta(n_sim_small,alpha,beta)
# curve of posterior with 20 observations
a_20 = alpha + sum(sample_small)
b_20 = beta + length(sample_small)-sum(sample_small)
curve(dbeta(x, a_20, b_20), lwd=1.5, ylim=c(0,30),
      xlim=c(0,0.8), main="curve of posterior with different observations",
      ylab="", col="dark red")

n_sim_medium = 200
# generation of sample from Beta with 200 observation
sample_medium = rbeta(n_sim_medium,alpha,beta)
# curve of posterior with high observations
```

```

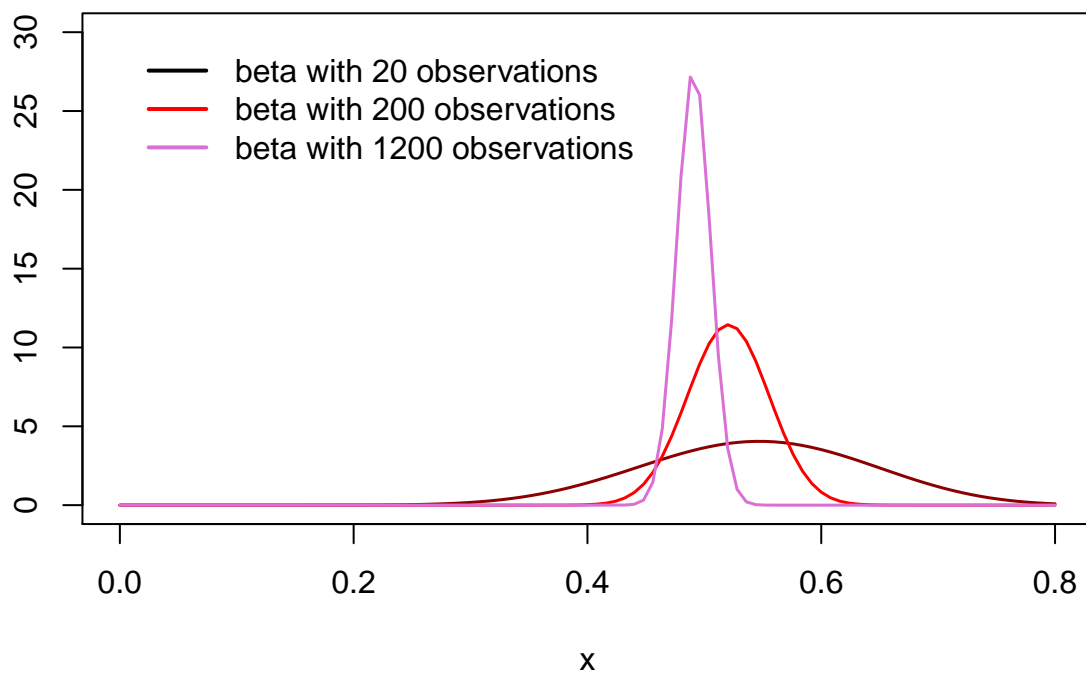
a_200 = alpha + sum(sample_medium)
b_200 = beta + length(sample_medium)-sum(sample_medium)
curve(dbeta(x, a_200, b_200), add=T, col="red", lwd=1.5)

n_sim_big = 1200
# generation of sample from Beta with 1200 observation
sample_big = rbeta(n_sim_big,alpha,beta)
# curve of posterior with high observations
a_1200 = alpha + sum(sample_big)
b_1200 = beta + length(sample_big)-sum(sample_big)
curve(dbeta(x, a_1200, b_1200), add=T, col="orchid", lwd=1.5)

legend(x=0,y=30, c("beta with 20 observations", "beta with 200 observations",
                  "beta with 1200 observations"), lty=c(1,1), lwd=c(2,2),
        col=c("black", "red","orchid"), bty="n")

```

curve of posterior with different observations



2e answer)

```

set.seed(123)

# compute the likelihood function of a bernoulli
L = function(x, theta){
  n = length(x)
  out = theta^(sum(x))*(1-theta)^(n-sum(x))
}

```

```

    return(out)
}

# simulate from a bernoulli
theta= 0.5
sample = rbinom(1000, 1, theta)
n = length(sample)

# find the maximum value of the likelihood function
k = L(sample, mean(sample))

# chose the parameters of the prior and compute it
alpha = 3
beta = 3
prior = function(x) return(dbeta(x, alpha, beta))


posterior = function(x) return(L(sample, x)*prior(x))

bounding_function = function(x) return(k*prior(x))

curve(posterior(x), col="dark red", lwd=2, ylim=c(0,3e-301),
      main = "target density and bounding function")
curve(bounding_function(x), lwd=2, add=T, col="red")
legend(x=0.60, y=3e-301, c("target density", "bounding function"),
      lty=c(1,1), lwd=c(2,2), col=c("dark red", "red"), bty="n")

```

target density and bounding function



```

# verify that theta has the target function as distribution


# generate  $Y_i$  where  $i=1,\dots,10000$  from  $q(y) \sim \text{Beta}(3,3)$ 
nsim = 10000
y = rbeta(nsim,alpha, beta)
U = runif(nsim,0,1)

# acceptance condition
theta_final = y[U <= L(sample, y)/k]
num_accepted=length(theta_final)
accept_probability=num_accepted/nsim

hist(theta_final, prob=T, main="empirical and target distribution",
      xlab="theta", ylim=c(0,30), ylab="", breaks = 50)
a = alpha + sum(sample)
b = beta + n - sum(sample)
curve(dbeta(x, a, b), add=T, col="red", lwd=2)

```

empirical and target distribution



Part 3

Acceptance/Rejection method - part II

3 answer)

We can simulate from a standard Normal distribution using pseudo-random deviates from a standard Cauchy and the A-R algorithm.

About the Acceptance/Rejection method, as seen before, to simulate we need:

- 1) target distribution written $f_X(x)$ from which we can compute it and in this case $f_X(x) \sim N(0, 1)$
- 2) a candidate distribution q that is easy to simulate (in this case $q \sim Cauchy(0, 1)$), so we can simulate Y_1, \dots, Y_n from candidate distribution $Y \sim q$

Our purpose corresponds to write the expression the corresponding acceptance probability of and evaluate it numerically by MC approximation.

About the choice of k , we know that the acceptance probability is high if the k value is smaller than possible, but at the same time k must satisfy the dominated condition, so the minimum value that k can assume is the maximum of function $\frac{f_X(x)}{q}$

The acceptance probability can be computed as $\frac{1}{k}$ or by Monte Carlo approximation that is equal to the ratio between the number of values accepted and the number of values generated:

$$\mathbb{E}(Y^A) \simeq \frac{1}{n} \sum_{i=1}^n Y^A$$

```
normal = function(x) return (dnorm(x,0,1))
cauchy = function(x) return(dcauchy(x,0,1))

d = function(x) return(normal(x)/cauchy(x))
curve(d(x),xlim=c(-5,5), ylim=c(0,2), ylab="", main="ratio between normal and cauchy")
opt = optimize(d, interval=c(-5,5), maximum=TRUE)
# maximum of the ratio between a normal and cauchy
maximum=opt$maximum
maximum


## [1] -0.9999966

# value of the maximum of the ratio between a normal and cauchy
k2=opt$objective
k2

## [1] 1.520347


#curve(d(x), xlim=c(-5,5), ylim=c(0,2), ylab="ratio between normal and cauchy")
abline(h=k2, col='red')
```

ratio between normal and cauchy



```
bounding_function2 = function(x) return(k2*cauchy(x))
curve(normal(x), xlim = c(-4,4), ylim=c(0,0.5), lwd=2,
      main = "target density and bounding function")
curve(bounding_function2, add=T, col="red", lwd=2)
legend(x=-4, y=0.5, c("target density","bounding function"),
      lwd=c(2,2), col=c("dark red", "red"), bty = "n")
```

target density and bounding function



```
# generate  $Y_i$  where  $i=1, \dots, 10000$  from  $q(y) \sim \text{cauchy}(0,1)$ 
nsim=10000
y2 = rcauchy(nsim,0,1)
U2 = runif(nsim,0,1)
```

```
x = y2[U2<=dnorm(y2,0,1)/(k2*dcauchy(y2,0,1))]
num_accepted2 = length(x)
```

```
# empirical acceptance probability
accept_probability2 = num_accepted2/nsim
accept_probability2
```

```
## [1] 0.6533
```

```
#theoretical acceptance probability ( $1/k2$ )
th_accept_probability2 = 1/k2
th_accept_probability2
```

```
## [1] 0.6577446
```

```
# comparison between empirical distribution and theoretical distribution
hist(x, prob=T, ylim=c(0,0.45), breaks = 50, ylab="",
     main="empirical and target distribution")
curve(normal(x), add=T, col="red", lwd=2)
```


empirical and target distribution

